cādence®

Fusion G3 Neural Network Library

Programmer's Guide - API

Contents

Nomenclature

Below naming convension is used while defining the APIs for kernels.

Naming convension	Description
f32	Single precision floating point
8	Signed 8-bit
8u	Unsigned 8-bit
16	Signed 16-bit
16u	Unsigned 16-bit
32	Signed 32-bit
32u	Unsigned 32-bit
Sym	Symmetric
Asym	Asymmetric
Axis	Dimension along which the computations will be performed
WORD32	32-bit signed integer
WORD16	16-bit signed half word
WORD8	Signed byte
UWORD32	32-bit unsigned integer
UWORD16	16-bit unsigned half word
UWORD8	Unsigned byte
FLOAT32	Single precision floating point

Document Change History

Version	Changes
1.0	■ Initial version

1. Introduction to the Fusion G3 NN library

The Fusion G3 Neural Network (NN) Library is an optimized implementation of various low-level NN kernels. The low-level NN kernels are the basic building blocks for operators and networks in neural network frameworks with a generic and simple interface.

Note   This version of the library supports Fusion G3 DSPs with the SP-VFPU (Single Precision Vector Floating Point Unit).

Note   This version of the Fusion G3 NN Library is tested with the xt-clang/xt-clang++ compilers using Xtensa Software Tools from RI-2022.10 release.

The Fusion G3 NN Library package includes the source code containing low-level kernel implementations.

This document covers information related to low level kernel APIs and information required to create the Fusion G3 NN Library. Section 2 provides details of low-level NN kernel APIs. Section 3 provides details of creating the Fusion G3 NN library.

1.1 Fusion G3 NN Library Specification

The current version of the Fusion G3 NN Library provides the following Fusion-optimized low-level kernel implementations.

1.1.1 Low-Level Kernels

- Activation kernels
- Basic operations kernels
    - o Quantize and dequantize kernels
    - o Basic vector math operators
    - o Broadcast kernels
- Normalization kernels
- Reorg kernels

These kernels support fixed point 8-bit, 16-bit, 32-bit, single precision floating point (float32/f32) data types for input and output. Please note that not all the kernels support all the datatypes specified here. The details of what datatypes are supported by each of the kernel is specified in Section 3. float32 is IEEE-754 compliant data types.

1.1.2 Support for Executorch Operators

The Fusion G3 NN Library low-level kernels can be used to implement the following operators of Executorch. 4-bit, 8-bit, 16-bit, 32-bit represents signed or unsigned datatypes. Kernels shown as supporting these datatypes might not support both signed and unsigned representations. Please refer Section 3 for details of the datatypes supported for each of the kernel.

No.	Operator	Float32 Datatype Support	32-bit	16-bit	8-bit	4-bit
1	Add	Yes	Yes	No	No	No
2	sub	Yes	Yes	No	No	No
3	mul	Yes	Yes	No	No	No
4	div	Yes	Yes	No	No	No
5	quantize	No	No	Yes	Yes	yes
6	dequantize	No	No	Yes	Yes	Yes
7	Softmax	Yes	No	No	No	No
8	Layernorm	Yes	No	No	No	No
9	Permute	No	Yes	Yes	Yes	No
10	Exp	Yes	No	No	No	No
11	cat	No	Yes	Yes	Yes	No
12	Slice	No	Yes	Yes	Yes	No
13	Clamp	Yes	No	Yes	Yes	No
14	sigmoid	Yes	No	No	No	No
15	sqrt	Yes	No	No	No	No
16	rsqrt	Yes	No	No	No	No
17	Tanh	Yes	No	No	No	No
18	Mean	Yes	No	No	No	No
19	Where	Yes	No	No	No	No
20	lt	Yes	No	No	No	No
21	Transpose	No	Yes	Yes	Yes	No

2. Fusion G3 NN Library – Low-Level Kernels

This section explains the APIs of low-level kernels which will be part of the NN library on Fusion G3 DSP. These kernels will be developed using Executorch as reference implementation. All the low-level kernels have a generic and simple interface.

2.1 Activation Kernels

2.1.1 Sigmoid

Description

The Sigmoid kernels perform the sigmoid operation on input vector `x` and give output vector as `y = sigmoid(x)`. Both the input and output vectors have size `vec_length`.

Function variants available are `xa_nn_sigmoid_[p]_[q]`, where:

> `[p]`: Input precision in bits

> `[q]`: Output precision in bits

Precision

Type	Description
f32_f32	float32 input, float32 output

Algorithm

$$y_n = \frac{1}{1 + \exp{(-x_n)}}, \qquad n = 0, \ldots, vec\_length - 1$$

Prototype

```
WORD32 xa_nn_sigmoid_f32_f32
(FLOAT32 * p_out,          const FLOAT32 * p_inp,          WORD32 vec_length);
```

Arguments

Type	Name	Size	Description
Input			
const	p_inp	vec_length	Input vector

FLOAT32 *			
WORD32	vec_length	1	Length of input vector
Output			
FLOAT32 *	p_out	vec_length	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp, p_out	Must not overlap
	Cannot be NULL
vec_length	Greater than 0

2.1.2 Tanh

Description

The Tanh kernels perform the hyperbolic tangent operation on input vector `x` and give output vector as `y` = `tanh(x)`. Both the input and output vectors have size `vec_length`.

Function variants available are `xa_nn_tanh_[p]_[q]`, where:

[p]: Input precision in bits

[q]: Output precision in bits

Precision

Type	Description
f32_f32	float32 input, float32 output

Algorithm

$$y_n = \tanh(x_n), \qquad n = 0, \ldots, vec\_length - 1$$

Prototype

```
WORD32 xa_nn_tanh_f32_f32
(FLOAT32 * p_out,          const FLOAT32 * p_inp,          WORD32 vec_length);
```

Arguments

Type	Name	Size	Description
Input			
const	p_inp	vec_length	Input vector

Type	Name	Size	Description
FLOAT32 *			
WORD32	vec_length	1	Length of input vector
Output			
FLOAT32 *	p_out	vec_length	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp, p_out	Must not overlap
	Cannot be NULL
vec_length	Greater than 0

2.1.3 Softmax

Description

The Softmax kernels compute the softmax (normalized exponential function) of input vector `x` and give output vector as `y = softmax(x)`. Both the input and output vectors have the same dimensions and size.

Function variants available are `xa_nn_softmax_[p]_[q]`, where:

[p]: Input precision in bits

[q]: Output precision in bits

Precision

Type	Description
f32_f32	float32 input, float32 output

Algorithm

$$y_n = \frac{\exp(x_n - max)}{\sum_k \exp(x_k - max)}, \qquad axis = NULL, \qquad n = \prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\text{-}shape[i] \ , \quad k = \prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$$

When axis is not NULL, below is the algorithm used to compute softmax.

$for\ j=0\ to\ \ \prod_{i=0}^{i=axis-1} p\_inp\_shape[i]$

$\quad outer\_stride = j * \prod_{i=axis}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$

$\quad for\ k=0\ to\ \prod_{i=axis+1}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$

$\qquad max\_value = 0;$

$\qquad accum = 0;$

$\qquad \textbf{end} = \prod_{i=axis}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$

$\qquad \textbf{stride} = \prod_{i=axis+1}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$

$\qquad for\ m = (outer\_stride + k)\ to\ \textbf{end}\ with\ a\ stride\ of\ \textbf{stride}$

$\qquad\qquad max\_value = max(max\_value, input[m])$

$\qquad for\ m = (outer\_stride + k)\ to\ \textbf{end}\ with\ a\ stride\ of\ \textbf{stride}$

$\qquad\qquad exp[m] = exp(input[m]\text{-}max\_value)$

$\qquad\qquad accum += exp[m]$

$\qquad for\ m = (outer\_stride + k)\ to\ \textbf{end}\ with\ a\ stride\ of\ \textbf{stride}$

$\qquad\qquad out[i] = exp[m]/accum$

For ex: For an input dimension of [20][12][4][3][3] with the axis set to 2, softmax is computed for groups of 4 elements each, where the dimension along axis 2 is 4. The softmax operation is performed separately for each group of elements in the last three dimensions (4x3x3) as illustrated in the figure below. This process is repeated across the top two dimensions of the input, which are dimensions 20 and 12.

Input vector

Output vector
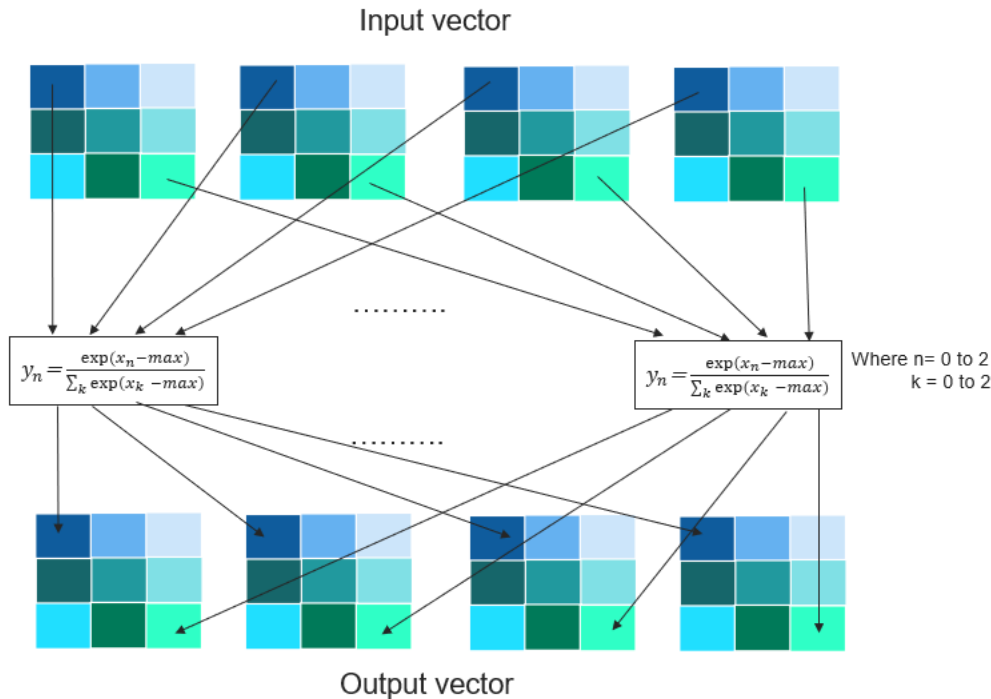
Prototype

```
WORD32 xa_nn_softmax_f32_f32
(FLOAT32 * p_out,          const FLOAT32 * p_inp,          const WORD32 * p_inp_shape,
WORD32 num_inp_dims,       WORD32 *p_axis);
```

Arguments

Type	Name	Size	Description
Input			
const FLOAT32 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input vector
const WORD32 *	p_inp_shape	num_inp_dims	Input shape
WORD32	num_inp_dims	1	Number of dimensions in the input
WORD32 *	p_axis	1	A dimension along which Softmax will be computed
Output			
FLOAT32 *	p_out	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Output vector.

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp, p_out	Must not overlap
	Cannot be NULL
p_inp_shape	Cannot be NULL
	Dimension values greater than 0
num_inp_dims	Greater than 0
*p_axis	[0, num_inp_dims)

2.2   Basic Operations and Miscellaneous Kernels

2.2.1 Elementwise Quantize Kernels

Description

The Elementwise Quantize kernels perform the quantization operation of the input vector elements to get the output vector. The kernels are developed in reference to the Quantize operator implementation in the Executorch. The Quantize kernels support both symmetric and asymmetric quantization for 4-bit (left

justified in byte), 8-bit and 16-bit datatypes (includes signed and unsigned datatypes). Function variants available are `xa_nn_elm_quantize_[p]_[q]`, where:

[p]: Input precision

[q]: Output precision

Precision

Type	Description
f32_asym4	Asymmetric quantization - single precision float input, signed 4-bit output (left justified in byte)
f32_asym8	Asymmetric quantization - single precision float input, signed 8-bit output
f32_asym16	Asymmetric quantization - single precision float input, signed 16-bit output
f32_sym4	Symmetric quantization - single precision float input, signed 4-bit output (left justified in byte)
f32_sym8	Symmetric quantization - single precision float input, signed 8-bit output
f32_sym16	Symmetric quantization - single precision float input, signed 16-bit output
f32_asym4u	Asymmetric quantization - single precision float input, unsigned 4-bit output (left justified in byte)
f32_asym8u	Asymmetric quantization - single precision float input, unsigned 8-bit output
f32_asym16u	Asymmetric quantization - single precision float input, unsigned 16-bit output
f32_sym4u	Symmetric quantization - single precision float input, unsigned 4-bit output (left justified in byte)
f32_sym8u	Symmetric quantization - single precision float input, unsigned 8-bit output
f32_sym16u	Symmetric quantization - single precision float input, unsigned 16-bit output

Algorithm

Tensor quantization

Asymmetric quantization

$for\ itr = 0{:}(num\_elm\text{-}1)$

$out\_value = (int)(p\_inp[itr]\ /\ out\_scale)\ +\ out\_zero\_bias\ ;$

$p_{out[itr]} = min(max(out\_value,\ quant\_min), quant\_max)\ ;$

Symmetric quantization

$for\ itr = 0{:}(num\_elm\text{-}1)$

$out\_value = (int)(p\_inp[itr]\ /\ out\_scale)$

$p\_out[itr] = min(max(out\_value,\ quant\_min), quant\_max)$

In both the above cases, for 4-bit quantized output

$for\ itr = 0{:}(num\_elm\text{-}1)$

$p\_out[itr] = p\_out[itr] << 4$

For channel/axis specific quantization – assuming the axis to be 2

Asymmetric quantization

$for\ itr = 0{:}(p\_inp\_shape[axis])$

$out\_value = (int)(p\_inp[d0][d1][itr][d3][d4]\ /\ out\_scale[iter])\ +\ out\_zero\_bias[iter]\ ;$

$p\_out[d0][d1][itr][d3][d4] = min(max(out\_value,\ quant\_min), quant\_max);$

Symmetric quantization

$for\ itr = 0{:}(p\_inp\_shape[axis])$

$out\_value = (int)(p\_inp[d0][d1][itr][d3][d4]\ /\ out\_scale[iter])\ ;$

$p\_out[d0][d1][itr][d3][d4] = min(max(out\_value,\ quant\_min), quant\_max);$

In both the above cases, for 4-bit quantized output

$p\_out[d0][d1][d2][d3][d4] = p\_out[d0][d1][d2][d3][d4] << 4\ ;$

Where

$d0 = 0\ to\ p\_inp\_shape[0]\text{-}1$
$d1 = 0\ to\ p\_inp\_shape[1]\text{-}1$
$d2 = 0\ to\ p\_inp\_shape[2]\text{-}1$
$d3 = 0\ to\ p\_inp\_shape[3]\text{-}1$
$d4 = 0\ to\ p\_inp\_shape[4]\text{-}1$

Both tensor quantization and channel based quantization is performed using the same kernel function.

Prototype

```
WORD32 xa_nn_elm_quantize_f32_asym4
(WORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 *p_out_zero_bias,            WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_asym8
(WORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 *p_out_zero_bias,            WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_asym16
(WORD16 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 *p_out_zero_bias,            WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym4
(WORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 quant_min,                   WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym8
(WORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 quant_min,                   WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym16
(WORD16 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_out_scale,              WORD32 quant_min,                   WORD32 quant_max);

WORD32 xa_nn_elm_quantize_f32_asym4u
(UWORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
```

```
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 *p_out_zero_bias,                  WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_asym8u
(UWORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 *p_out_zero_bias,                  WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_asym16u
(UWORD16 *__restrict__ p_out,       const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 *p_out_zero_bias,                  WORD32 quant_min,
WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym4u
(UWORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 quant_min,                        WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym8u
(UWORD8 *__restrict__ p_out,        const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 quant_min,                        WORD32 quant_max);
WORD32 xa_nn_elm_quantize_f32_sym16u
(UWORD16 *__restrict__ p_out,       const FLOAT32 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,    WORD32 num_inp_dims,                      WORD32 *p_axis,
FLOAT32 *p_out_scale,               WORD32 quant_min,                        WORD32 quant_max);
```

Arguments

Type	Name	Size	Description
Input			
const FLOAT32 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input vector
Const WORD32 *const	P_inp_shape	num_inp_dims	Shape of the input vector
WORD32	num_inp_dims	1	Number of input dimensions
WORD32 *	p_axis	1	A dimension along which quantization will be computed
FLOAT32 *	p_out_scale	1 Or P_inp_shape[*p_axis]	Scale of output
WORD32 *	p_out_zero_bias	1 Or P_inp_shape[*p_axis]	Zero offset of output.
WORD32	quant_min	1	Minimum value used to limit the output
WORD32	quant_max	1	Maximum value used to limit the output
Output			
WORD8 * WORD16 * UWORD8 * UWORD16 *	p_out	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element of datatype)-byte boundary
	Cannot be NULL
	Must not overlap
num_inp_dims	Greater than 0 and less than or equal 5
*p_out_scale	Not equal to zero and finite single precision float value
*p_out_zero_bias	{-8,.......,7} for out type signed 4-bit
	{-128,...., 127} for out type signed 8-bit
	{-32768...., 32767} for out type signed 16-bit
	{0,.......,15} for out type unsigned 4-bit
	{0,...., 255} for out type unsigned 8-bit
	{0...., 65535} for out type unsigned 16-bit
*p_axis	[0, num_inp_dims)

2.2.2 Elementwise Dequantize Kernels

Description

The Elementwise Dequantize kernels perform the dequantization operation of the input vector elements to get the output vector. The dequantize kernels support both symmetric and asymmetric dequantization for 4-bit (left justified in byte), 8-bit and 16-bit datatypes (includes signed and unsigned datatypes).

Function variants available are `xa_nn_elm_dequantize_[p]_[q]`, where:

[p]: Input precision

[q]: Output precision

Precision

Type	Description
asym4_f32	Asymmetric dequantization - signed 4-bit input (left justified in byte), single precision float output
asym8_f32	Asymmetric dequantization - signed 8-bit input, single precision float output
asym16_f32	Asymmetric dequantization - signed 16-bit input, single precision float output
sym4_f32	Symmetric dequantization - signed 4-bit input (left justified in byte), single precision float output
sym8_f32	Symmetric dequantization - signed 8-bit input, single precision float output
sym16_f32	Symmetric dequantization - signed 16-bit input, single precision float output
asym4u_f32	Asymmetric dequantization - unsigned 4-bit input (left justified in byte), single precision float output
asym8u_f32	Asymmetric dequantization - unsigned 8-bit

	input, single precision float input
asym16u_f32	Asymmetric dequantization - unsigned 16-bit input, single precision float output
sym4u_f32	Symmetric dequantization - unsigned 4-bit input (left justified in byte), single precision float output
sym8u_f32	Symmetric dequantization - unsigned 8-bit input, single precision float output
sym16u_f32	Symmetric dequantization - unsigned 16-bit input, single precision float output

Algorithm

Tensor dequantization

In both the below cases, for 4-bit quantized input

$for\ itr = 0:(num\_elm\text{-}1)$

$p\_inp[itr] = p\_inp[itr] >> 4;$

Asymmetric dequantization

$for\ itr = 0:(num\_elm\text{-}1)$

$p\_out[itr] = (float)((p\_inp[itr] - inp\_zero\_bias) * inp\_scale))$

Symmetric dequantization

$for\ itr = 0:(num\_elm\text{-}1)$

$p\_out[itr] = (float)((p\_inp[itr]) * inp\_scale))$

For channel/axis specific dequantization – assuming the axis to be 2

In both the below cases, for 4-bit quantized input

$p\_inp[d0][d1][d2][d3][d4] = p\_inp[d0][d1][d2][d3][d4] >> 4;$

Asymmetric quantization

$for\ itr = 0:(p\_inp\_shape[axis])$

$p\_out[d0][d1][itr][d3][d4] = (float)((p\_inp[d0][d1][itr][d3][d4] - inp\_zero\_bias[itr]) * inp\_scale[itr]));$

Symmetric quantization

$for\ itr = 0:(p\_inp\_shape[axis])$

$p\_out[d0][d1][itr][d3][d4] = (float)(p\_inp[d0][d1][itr][d3][d4] * inp\_scale[itr]);$

Where
  d0 = 0 to p_inp_shape[0]-1
  d1 = 0 to p_inp_shape[1]-1
  d2 = 0 to p_inp_shape[2]-1
  d3 = 0 to p_inp_shape[3]-1
  d4 = 0 to p_inp_shape[4]-1

Note: Both tensor quantization and channel based quantization is performed using the same kernel function.

Prototype

```
WORD32 xa_nn_elm_dequantize_asym4_f32
(FLOAT32 *__restrict__ p_out,      const WORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_asym8_f32
(FLOAT32 *__restrict__ p_out,      const WORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_asym16_f32
(FLOAT32 *__restrict__ p_out,      const WORD16 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym4_f32
(FLOAT32 *__restrict__ p_out,      const WORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym8_f32
(FLOAT32 *__restrict__ p_out,      const WORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym16_f32
(FLOAT32 *__restrict__ p_out,      const WORD16 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);


WORD32 xa_nn_elm_dequantize_asym4u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_asym8u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_asym16u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD16 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
WORD32 *p_inp_zero_bias,           FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym4u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym8u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD8 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);
WORD32 xa_nn_elm_dequantize_sym16u_f32
(FLOAT32 *__restrict__ p_out,      const UWORD16 *__restrict__ p_inp,
const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,                WORD32 *p_axis,
FLOAT32 *p_inp_scale);
```

Arguments

Type	Name	Size	Description
Input			

const WORD8 * WORD16 * UWORD8 * UWORD16 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\text{-}shape[i]$	Input vector
Const WORD32 *contest	p_inp_shape	num_inp_dims	Shape of the input vector
WORD32	num_inp_dims	1	Number of input dimensions
WORD32 *	p_inp_zero_bias	1 Or P_inp_shape[*p_axis]	Zero offset of input
FLOAT32 *	p_inp_scale	1 Or P_inp_shape[*p_axis]	Input scale
WORD32 *	P_axis	1	A dimension along which dequantization will be computed
Output			
FLOAT32 *	p_out	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\text{-}shape[i]$	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Must not overlap
num_inp_dims	Greater than 0 and less than or equal 5
*p_inp_zero_bias	{-8,……..,7} for out type signed 4-bit {-128,…., 127} for out type signed 8-bit {-32768…., 32767} for out type signed 16-bit {0,……..,15} for out type unsigned 4-bit {0,…., 255} for out type unsigned 8-bit {0…., 65535} for out type unsigned 16-bit
*p_axis	[0, num_inp_dims)
*p_inp_scale	Finite single precision float value

2.2.3 Basic Vector math operation Kernels

Description

The Basic kernels perform basic elementwise operations on one or two input vectors x and y to get output vector z. The supported operations are: add, subtract, multiply, div, exp, clamp, where, square-root and inverse square-root. The supported precisions are: float32, int32, int16, int8, uint8.

Function variants available are xa_nn_elm_[o]_[p]_[q], where:

[o]: Operations: add, add_scalar, sub, sub_scalar, mul, mul_scalar, div, div_scalar, exp, clamp, clamp_scalar, where, sqrt, rsqrt

[p]: Input Precision in bits- input1xinput2 or input1

[q]: Output Precision in bits

Precision

Type	Description
f32xf32_f32	2 float32 inputs, float32 output
f32_f32	float32 input, float32 output
8_8	signed 8-bit input, signed 8-bit output
8u_8u	unsigned 8-bit input, unsigned 8-bit output
16_16	signed 16-bit input, signed 16-bit output
32x32_32	2 32-bit input, 32-bit output

Algorithm

$$elm\_add \quad : \quad z_n = x_n + \alpha * y_n\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_add\_scalar: \quad z_n = x_n + \alpha * y\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_sub \quad : \quad z_n = x_n - \alpha * y_n\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_sub\_scalar: \quad z_n = x_n - \alpha * y\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_mul \quad : \quad z_n = x_n * y_n\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_mul\_scalar: \quad z_n = x_n * y\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_div \quad : \quad z_n = x_n / y_n\,, \qquad n = 0\,....,num\_elm - 1 \qquad mode=0$$
$$\qquad : \quad z_n = truncate(x_n / y_n)\,, \quad n = 0\,....,num\_elm - 1, \quad mode=1$$
$$\qquad : \quad z_n = floor(x_n / y_n)\,, \qquad n = 0\,....,num\_elm - 1, \quad mode=2$$
$$elm\_div\_scalar: \quad z_n = x_n / y\,, \qquad n = 0\,....,num\_elm - 1, \quad mode=0$$
$$\qquad : \quad z_n = truncate(x_n / y)\,, \quad n = 0\,....,num\_elm - 1, \quad mode=1$$
$$\qquad : \quad z_n = floor(x_n / y)\,, \qquad n = 0\,....,num\_elm - 1, \quad mode=2$$
$$elm\_exp: \quad z_n = exp(x_n)\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_clamp: \quad z_n = \min(max(x_n, xmin_n), xmax_n)\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_clamp\_scalar: \quad z_n = \min(max(x_n, xmin), xmax) \qquad n = 0\,....,num\_elm - 1$$
$$elm\_where: \quad z_n = (a_n ? x_n : y_n)\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_sqrt: \quad z_n = \sqrt{x_n}\,, \qquad n = 0\,....,num\_elm - 1$$
$$elm\_rsqrt: \quad z_n = 1 \div \sqrt{x_n}\,, \qquad n = 0\,....,num\_elm - 1$$

x_n represents first input, y_n represents second input, α represents a scale value, a_n represents the condition on which any one of the inputs is selected in "where" operator, $xmin_n$ and $xmax_n$ represents array of minimum and maximum values used to limit the input in "clamp" operator. All the variables without a subscript "n" has the same meaning as above but they are scalar values.

z_n represents output.

Prototype

```
WORD32 xa_nn_elm_add_f32xf32_f32
(FLOAT32 * p_out,          const  FLOAT32 * p_inp1,     const  FLOAT32 * p_inp2,
FLOAT32 alpha,             WORD32 num_elm);
WORD32 xa_nn_elm_add_scalar_f32xf32_f32
```

```
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32  inp2,
FLOAT32 alpha,                   WORD32 num_elm);
WORD32 xa_nn_elm_add_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 * p_inp2,
WORD32 alpha,                    WORD32 num_elm);
WORD32 xa_nn_elm_add_scalar_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 inp2,
WORD32 alpha,                    WORD32 num_elm);

WORD32 xa_nn_elm_sub_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32 * p_inp2,
FLOAT32 alpha,                   WORD32 num_elm);
WORD32 xa_nn_elm_sub_scalar_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32  inp2,
FLOAT32 alpha,                   WORD32 num_elm);
WORD32 xa_nn_elm_sub_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 * p_inp2,
WORD32 alpha,                    WORD32 num_elm);
WORD32 xa_nn_elm_sub_scalar_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 inp2,
WORD32 alpha,                    WORD32 num_elm);

WORD32 xa_nn_elm_mul_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32 * p_inp2,
WORD32 num_elm);
WORD32 xa_nn_elm_mul_scalar_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32  inp2,
WORD32 num_elm);
WORD32 xa_nn_elm_mul_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 * p_inp2,
WORD32 num_elm);
WORD32 xa_nn_elm_mul_scalar_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 inp2,
WORD32 num_elm);


WORD32 xa_nn_elm_div_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32 * p_inp2,
WORD32 mode,                     WORD32 num_elm);
WORD32 xa_nn_elm_div_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 * p_inp2,
WORD32 mode,                     WORD32 num_elm);
WORD32 xa_nn_elm_div_32x32_f32
(FLOAT32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 * p_inp2,
WORD32 num_elm);


WORD32 xa_nn_elm_div_scalar_f32xf32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp1,     const  FLOAT32 inp2,
WORD32 mode,                     WORD32 num_elm);
WORD32 xa_nn_elm_div_scalar_32x32_32
(WORD32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 inp2,
WORD32 mode,                     WORD32 num_elm);
WORD32 xa_nn_elm_div_scalar_32x32_f32
(FLOAT32 * p_out,                 const  WORD32 * p_inp1,      const  WORD32 inp2,
WORD32 num_elm);


WORD32 xa_nn_elm_exp_f32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp,      WORD32 num_elm);

WORD32 xa_nn_elm_clamp_f32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp,      const  FLOAT32 * p_min,
const  FLOAT32 * p_max,          WORD32 num_elm);
WORD32 xa_nn_elm_clamp_scalar_f32_f32
(FLOAT32 * p_out,                const  FLOAT32 * p_inp,      const  FLOAT32 min,
const  FLOAT32 max,              WORD32 num_elm);
WORD32 xa_nn_elm_clamp_16_16
```

```
(WORD16 * p_out,            const  WORD16 * p_inp,         const  WORD16 * p_min,
const  WORD16 * p_max,        WORD32 num_elm);
WORD32 xa_nn_elm_clamp_scalar_16_16
(WORD16 * p_out,            const  WORD16 * p_inp,         const  WORD16 min,
const  WORD16 max,           WORD32 num_elm);
WORD32 xa_nn_elm_clamp_8_8
(WORD8 * p_out,             const  WORD8 * p_inp,          const  WORD8 * p_min,
const  WORD8 * p_max,        WORD8 num_elm);
WORD32 xa_nn_elm_clamp_scalar_8_8
(WORD8 * p_out,             const  WORD8 * p_inp,          const  WORD8 min,
const  WORD8 max,            WORD8 num_elm);
WORD32 xa_nn_elm_clamp_8u_8u
(UWORD8 * p_out,            const  UWORD8 * p_inp,         const  UWORD8 * p_min,
const  UWORD8 * p_max,       WORD8 num_elm);
WORD32 xa_nn_elm_clamp_scalar_8u_8u
(UWORD8 * p_out,            const  UWORD8 * p_inp,         const  UWORD8 min,
const  UWORD8 max,           WORD8 num_elm);

WORD32 xa_nn_elm_where_f32xf32_f32
(FLOAT32 * p_out,           const  FLOAT32 * p_inp1,       const  FLOAT32 * p_inp2,
const  UWORD8 * p_cond,      WORD32 num_elm);

WORD32 xa_nn_elm_sqrt_f32_f32
(FLOAT32 * p_out,           const  FLOAT32 * p_inp,        WORD32 num_elm);

WORD32 xa_nn_elm_rsqrt_f32_f32
(FLOAT32 * p_out,           const  FLOAT32 * p_inp,        WORD32 num_elm);
```

Arguments

Type	Name	Size	Description
Input			
const UWORD8 * WORD8 * WORD16 * WORD32 * FLOAT32 *	p_inp1, p_inp,	num_elm	First input vector
const WORD32 * FLOAT32 *	p_inp2	num_elm	Second input vector
WORD32 FLOAT32	inp2	1	Second input which is scalar
WORD32	num_elm	1	Number of elements
WORD32 FLOAT32	alpha	1	Scale for the second operand in add and sub operators
UWORD8 * WORD8 * WORD16 * FLOAT32 *	p_min, min	num_elm	Minimum values vector
UWORD8 WORD8 WORD16 FLOAT32	min	1	Minimum value which is scalar
UWORD8 * WORD8 * WORD16 * FLOAT32 *	p_max	num_elm	Max values vector
UWORD8 WORD8 WORD16 FLOAT32	max	1	Maximum value which is scalar
UWORD8 *	p_cond	num_elm	Condition on which one of

Type	Name	Size	Description
			the input is selected in where operator
WORD32	mode	1	Type of division 0 – normal a/b 1 – truncate the result after division 2 – floor of the result after division
Output			
UWORD8 * WORD8 * WORD16 * WORD32 * FLOAT32 *	p_out	num_elm	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_inp1,p_inp2, p_inp, p_min, p_max, p_cond p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
p_out	Must not overlap with the input pointers
num_elm	Greater than 0
mode	0, 1, 2

2.2.4 Elementwise Comparison Kernels

Description

The Elementwise comparison kernels perform elementwise comparison operations on two input vectors x and y to get the output vector z. Currently, the supported operation is: less than (<). The output for the comparison kernels is a Boolean value that requires 1-byte space. The supported precisions are: f32.

Function variants available are xa_nn_elm_[o]_[p]_[q], where:

[o]: Operations: less

[p]: Input Precision in bits- input1xinput2

[q]: output - bool

Precision

Type	Description

`f32xf32_bool`	2 float32 inputs, Boolean(1-byte) output

Algorithm

elm_less: $\quad z_n = (x_n < y_n)\,, \qquad n = 0 \,....,num\_elm - 1$

elm_less_scalar: $\quad z_n = (x_n < y)\,, \qquad n = 0 \,....,num\_elm - 1$

x_n represents first input, $y_n \ or \ y$ represents second input.

z_n represents output.

Prototype

```
WORD32 xa_nn_elm_less_f32xf32_bool
(WORD8 * p_out,         const FLOAT32 * p_inp1,    const FLOAT32 * p_inp2,
 WORD32  num_elm);

WORD32 xa_nn_elm_less_scalar_f32xf32_bool
(WORD8 * p_out,         const FLOAT32 * p_inp1,    const FLOAT32 inp2,
 WORD32  num_elm);
```

Arguments

Type	Name	Size	Description
Input			
const FLOAT32 *	p_inp1	num_elm	First input vector
const FLOAT32 *	p_inp2	num_elm	Second input vector or scalar
FLOAT32	inp2	1	Second input as scalar
WORD32	num_elm	1	Number of elements
Output			
WORD8 *	p_out	num_elm	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_inp1,p_inp2, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
P_out	Must not overlap with the input pointers
num_elm	Greater than 0

2.2.5 Basic Kernels with 5D Broadcasting

Description

The Basic kernels with 5D broadcasting perform a broadcast operation and apply an operator on the inputs. The supported operators are: elementwise add, sub, mul, div, less, clamp and where.

Details of the broadcast operation can be found at [Executorch Broadcasting semantics](#).

These kernels support upto 5-dimensional input/output tensors. The 1/2/3/4-dimensional inputs can be scaled up to 5D within the kernel. Both inputs and output must have the same number of dimensions.

Tensors must also be broadcast compatible (that is, either their dimensions must match or be equal to 1) otherwise kernels return error.

Function variants available are `xa_nn_elm_[op]_broadcast_5D_[p]_[q]`, where:

> [op]: Operation: add, sub, mul, div, less, where, clamp
>
> `[p]`: Input Precision in bits- input1xinput2 or input1
>
> `[q]`: Output Precision in bits

Precision

Type	Description
f32_f32	Float32 input, float32 output
8_8	Signed 8-bit input, signed 8-bit output
8u_8u	Unsigned 8-bit input, unsigned 8-bit output
16_16	Signed 16-bit input, signed 16-bit output
32x32_32	2 32-bit input, 32-bit output
f32xf32_bool	2 float32 inputs, Boolean(1-byte) output
f32xf32_f32	2 f32 inputs,_f32 output

Algorithm

$$p\text{--}out[i_0][i_1]\dots[i_4] = [op](p\_inp1[i1_0][i1_1]\dots[i1_4], \; p\_inp2[i2_0][i2_1]\dots[i2_4])$$

Where,

- $i_n = \max(i1_n, i2_n)\,;\; n = [0,4]$

Ops are:

elm_add:	$z_n = x_i + \alpha * y_j$	
elm_sub:	$z_n = x_i - \alpha * y_j$	
elm_mul:	$z_n = x_i * y_j$	
elm_div:	$z_n = x_i / y_j,$	mode = 0
:	$z_n = truncate(x_i / y_j),$	mode = 1
:	$z_n = floor(x_i / y_j)\,,$	mode = 2

elm_less: $\qquad z_n = x_i < y_j$

elm_where: $\qquad z_n = (a_i ? \, x_j : y_k)$

elm_clamp: $\qquad z_n = \min \, (max(x_i, xmin_j), xmax_k)$

Prototypes

```
WORD32 xa_nn_elm_add_broadcast_5D_f32xf32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims,
 FLOAT32 alpha
);
WORD32 xa_nn_elm_add_broadcast_5D_32x32_32
(WORD32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const WORD32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims,
 WORD32 alpha
);
WORD32 xa_nn_elm_sub_broadcast_5D_f32xf32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims,
 FLOAT32 alpha
);
WORD32 xa_nn_elm_sub_broadcast_5D_32x32_32
(WORD32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const WORD32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims,
 WORD32 alpha
);
WORD32 xa_nn_elm_mul_broadcast_5D_f32xf32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims
```

```
);
WORD32 xa_nn_elm_mul_broadcast_5D_32x32_32
(WORD32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const WORD32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims
);


WORD32 xa_nn_elm_div_broadcast_5D_f32xf32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32          mode,
 WORD32       num_inp_dims
);
WORD32 xa_nn_elm_div_broadcast_5D_32x32_32
(WORD32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const WORD32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       mode,
 WORD32       num_inp_dims
);
WORD32 xa_nn_elm_div_broadcast_5D_32x32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const WORD32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims
);



WORD32 xa_nn_elm_less_broadcast_5D_f32xf32_bool
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 WORD32       num_inp_dims

);


WORD32 xa_nn_elm_where_broadcast_5D_f32xf32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
```

```
 const FLOAT32 * __restrict__ p_inp1,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_inp2,
 const WORD32 *const p_inp2_shape,
 const  UWORD8 * p_cond,
 const WORD32 *const p_cond_shape,
 WORD32        num_inp_dims
);


WORD32 xa_nn_elm_clamp_broadcast_5D_f32_f32
(FLOAT32 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const FLOAT32 * __restrict__ p_inp,
 const WORD32 *const p_inp1_shape,
 const FLOAT32 * __restrict__ p_min,
 const WORD32 *const p_min_shape,
 const  FLOAT32 * __restrict__ p_max,
 const WORD32 *const p_max_shape,
 WORD32        num_inp_dims
);


WORD32 xa_nn_elm_clamp_broadcast_5D_16_16
(WORD16 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD16 * __restrict__ p_inp,
 const WORD32 *const p_inp1_shape,
 const WORD16 * __restrict__ p_min,
 const WORD32 *const p_min_shape,
 const WORD16 * p_max,
 const WORD32 *const p_max_shape,
 WORD32        num_inp_dims
);


WORD32 xa_nn_elm_clamp_broadcast_5D_8_8
(WORD8 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const WORD8 * __restrict__ p_inp,
 const WORD32 *const p_inp1_shape,
 const WORD8 * __restrict__ p_min,
 const WORD32 *const p_min_shape,
 const WORD8 * p_max,
 const WORD32 *const p_max_shape,
 WORD32        num_inp_dims
);


WORD32 xa_nn_elm_clamp_broadcast_5D_8u_8u
(UWORD8 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const UWORD8 * __restrict__ p_inp,
 const WORD32 *const p_inp1_shape,
 const UWORD8 * __restrict__ p_min,
 const WORD32 *const p_min_shape,
 const UWORD8 * p_max,
 const WORD32 *const p_max_shape,
 WORD32        num_inp_dims
);
```

Arguments

Type	Name	Size	Description
Input			
const WORD8 *, const UWORD8 *, const WORD16 * const FLOAT32 *, const WORD32 *	p_inp1, p_inp	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_inp1\_shape[i]$	First input tensor
const FLOAT32 *, const WORD32 *	p_inp2	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_inp2\_shape[i]$	Second input tensor
const WORD8 *	p_cond	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_cond\_shape[i]$	Input tensor holding conditions used in where operator
const WORD8 *, const UWORD8 *, const WORD16 * const FLOAT32 *, const WORD32 *	p_max	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_max\_shape[i]$	Input tensor holding max values for Clamp operator
const WORD8 *, const UWORD8 *, const WORD16 * const FLOAT32 *,	p_min	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_min\_shape[i]$	Input tensor holding min values for Clamp operator
const WORD32 *const	p_out_shape	num_inp_dims	Shape of output (first dimension is outer most)
const WORD32 *const	p_inp1_shape p_inp_shape	num_inp_dims	Shape of first input (first dimension is outer most)
const WORD32 *const	p_inp2_shape	num_inp_dims	Shape of second input (first dimension is outer most)
const WORD32 *const	p_cond_shape	num_inp_dims	Shape of condition tensor input (first dimension is outer most)
const WORD32 *const	p_max_shape	num_inp_dims	Shape of max tensor input (first dimension is outer most)
const WORD32 *const	p_min_shape	num_inp_dims	Shape of min tensor input (first dimension is outer most)
WORD32	num_inp_dims	1	Number of input

Type	Name	Size	Description
			dimensions
WORD32	mode	1	Type of division 0 – normal a/b 1 – truncate the result after division 2 – floor of the result after division
Output			
UWORD8 * WORD8 * FLOAT32 * WORD16 * WORD32 *	p_out	$\displaystyle\prod_{i=0}^{i=num\_inp\_dims-1} p\_out\_shape[i]$	Output tensor

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp1, p_inp, p_inp2, p_cond, p_max, p_min p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
p_out	Must not overlap with the input pointers
p_out_shape, p_inp1_shape, p_inp2_shape, p_cond_shape, p_max_shape, p_min_shape	Cannot be NULL
	Aligned on 4-byte boundary
	Shapes must be broadcast compatible, that is, p_out_shape[i] must be max(ith shape of all the inputs]) p_inp1_shape[i] must be either equal to ith shape of other inputs or 1 p_inp2_shape[i] must be either equal to ith shape of other inputs or 1 p_cond_shape[i] must be either equal to ith shape of other inputs or 1 p_max_shape[i] must be either equal to ith shape of other inputs or 1 p_min_shape[i] must be either equal to ith shape of other inputs or 1

2.2.6 Mean

Description

The Mean kernel computes the mean of input vector `x` and gives output vector as `y = mean(x)`. The output vector dimension will be reduced while performing the operation.

Function variants available are `xa_nn_mean_[p]_[q]`, where:

> `[p]`: Input precision in bits

> `[q]`: Output precision in bits

Precision

There is single variant available:

Type	Description
f32_f32	float32 input, float32 output

Algorithm

$$y = \frac{\sum_k x_n}{K}, \qquad * p\_axis = NULL, \qquad k = \prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$$

When axis is not NULL, below is the algorithm used to compute mean.

out_idx = 0;

for j=0 to $\prod_{i=0}^{i=axis-1} p\_inp\_shape[i]$

    outer_stride = j* $\prod_{i=axis}^{i=num-inp-dims-1} p\_inp\_shape[i];$

    for k=0 to $\prod_{i=axis+1}^{i=num-inp-dims-1} p\_inp\_shape[i]$

        accum = 0;

        end $= \prod_{i=axis}^{i=num-inp-dims-1} p\_inp\_shape[i]$ *;*

        stride $= \prod_{i=axis+1}^{i=num-inp-dims-1} p\_inp\_shape[i];$

        for m = (outer_stride + k) to ***end*** *with a stride of* ***stride***

            accum = accum+input[m];

        out[out_idx] = accum/p_inp_shape[axis];

        out_idx = out_idx+1;

For ex: For an input dimension of [20][12][4][3][3] with the axis set to 2, mean is computed for groups of 4 elements each, where the dimension along axis 2 is 4. The mean is computed on a group of 4 elements each from the last 3 dimensions (4x3x3) as illustrated in the figure below. This process is repeated across the top two dimensions of the input, which are dimensions 20 and 12.

Input vector

$$y = \frac{\sum_k x_k}{K}$$

$$y = \frac{\sum_k x_k}{K}$$

Where k = 0 to 3
K = 4

Output vector
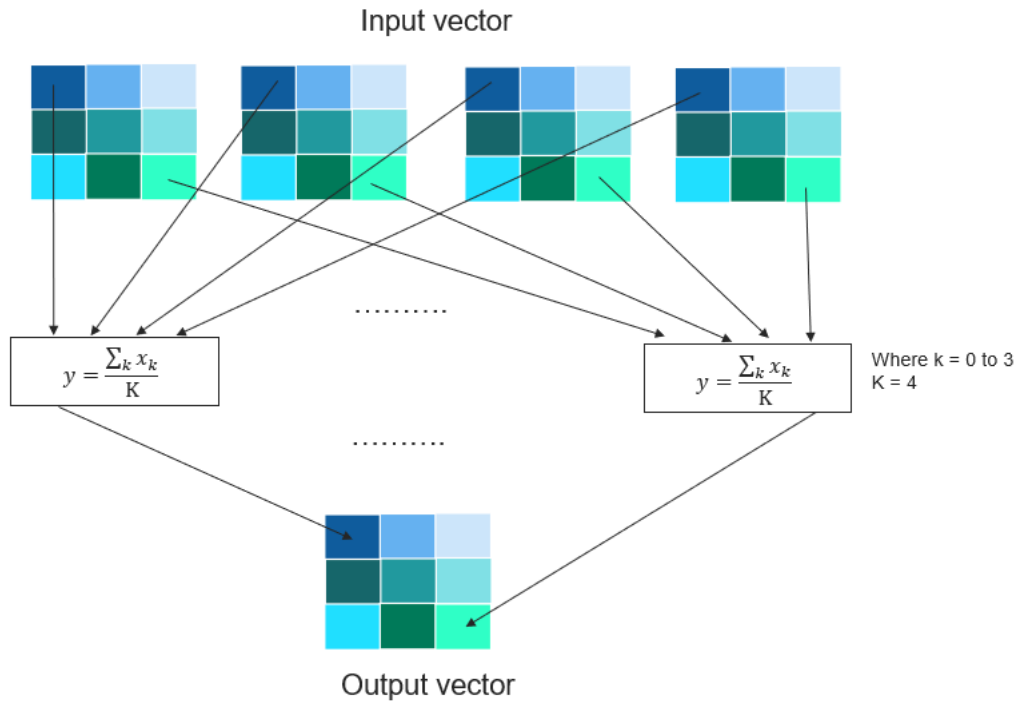
Prototype

```
WORD32 xa_nn_mean_f32_f32
(FLOAT32 * p_out,              const WORD32 * p_out_shape,         WORD32 num_out_dims,
const FLOAT32 * p_inp,         const WORD32 * p_inp_shape,    WORD32 num_inp_dims,
WORD32 *p_axis);
```

Arguments

Type	Name	Size	Description
Input			
const FLOAT32 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input vector
const WORD32 *	p_inp_shape	num_inp_dims	Input shape
const WORD32 *	p_out_shape	num_out_dims	Output shape
WORD32	num_inp_dims	1	Number of dimensions in the input
WORD32	num_out_dims	1	Number of dimensions in the output
WORD32	p_axis	1	A dimension along which mean will be computed
Output			

FLOAT32 *	p_out	$\prod_{i=0}^{i=num\_out\_dims-1} p\_out\_shape[i]$	Output vector.

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	must not overlap
	Cannot be NULL
p_inp_shape, p_out_shape	The dimension values should be greater than 0
	Aligned on 4-byte boundary
num_inp_dims num_out_dims	Greater than 0
Dim	Maximum value is num_inp_dims-1
*p_axis	[0, num_dim)

2.3   Normalization Kernels

2.3.1 Layer Normalization Kernel

Description

The Layer Normalization kernel applies Layer normalization on an input vector `x` to get output vector `z`.

Function variants available are `xa_nn_native_layer_norm_[p]_[q]`, where:

`[p]`: Input precision in bits

`[q]`: Output precision in bits

Precision

Type	Description
f32_f32	float32 input, float32 output

Algorithm

$$z_i[n] = \frac{x_i[n] - \mu_i}{\sigma_i + eps} * \; w[n] + b[n]$$

$for\ a\ given\ axis\ k\ and\ dimension\ size\ \text{num\_inp\_dims}$

$$i = \prod_{j=0}^{j=k-1} p\_inp\_shape[j]$$

$$n = \prod_{j=k}^{j=num\_inp\_dims-1} p\_inp\_shape[j]$$

$$\mu_i = mean(x_i[n])$$
$$\sigma_i = std(x_i[n])$$

$$z_n = \frac{x_n - \mu}{\sigma + eps} * w_n + b_n$$

when axis is 0

$$n = \prod_{j=0}^{j=num\_inp\_dims-1} p\_inp\_shape[j]$$

x_n *and* $x_i[n]$ represents input vector.

w_n *and* $w[n]$ represents weight vector.

b_n *and* $b[n]$ represents bias vector.

μ *and* μ_i represents mean.

σ *and* σ_i represents standard deviation.

z_n *and* $z_i[n]$ represents output vector.

For ex: For an input dimension of [20][12][4][3][3] and with the axis set to 2, layer normalization is applied to groups of 36 elements, where each group consists of a 4x3x3 block of elements. This operation is performed independently for each group, as illustrated in the figure below. The process is repeated across the top two dimensions of the input, namely dimensions 20 and 12..
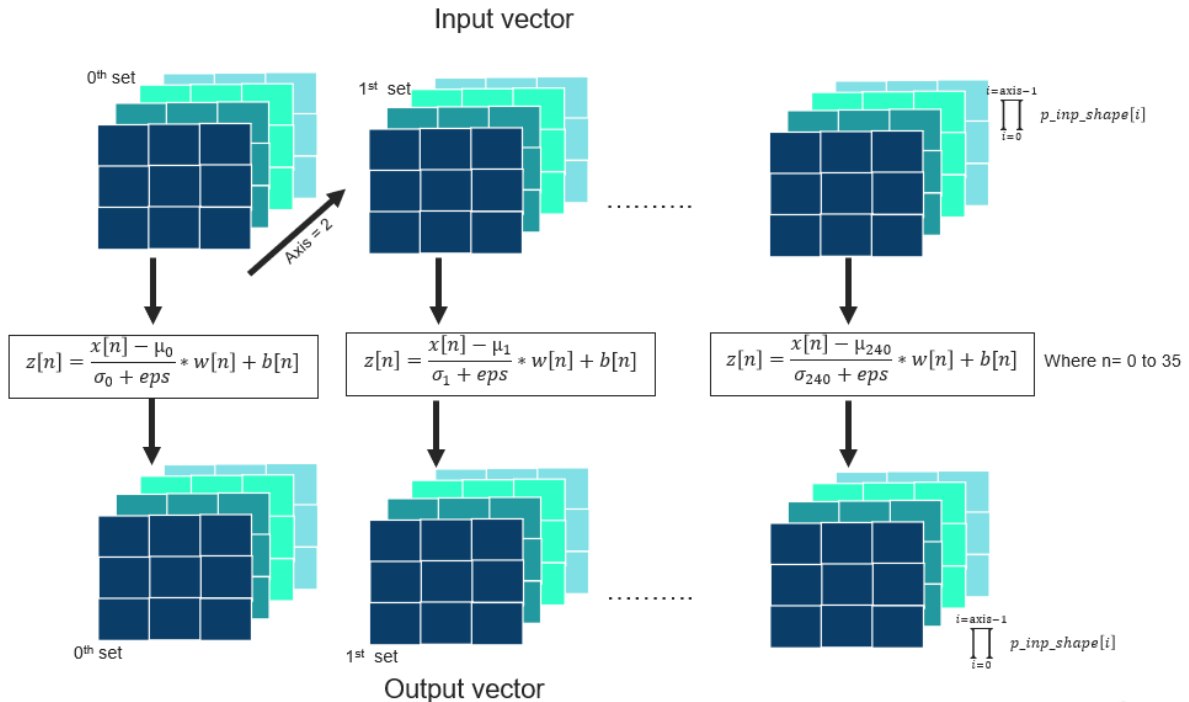
Prototype

```
WORD32 xa_nn_native_layer_norm_f32_f32
(FLOAT32 * p_out,                      FLOAT32 * p_mean,                  FLOAT32 * p_rstd,
 const FLOAT32 * p_inp,          const WORD32 *const p_inp_shape,   WORD32 num_inp_dims,
 WORD32 axis,                 const FLOAT32 * p_weight,        const FLOAT32 * p_bias,
 FLOAT32  eps);
```

Arguments

Type	Name	Size	Description
Input			
const FLOAT32 *	p_inp	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input vector
const FLOAT32 *	p_weight	$\prod\limits_{i=*p\_axis}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Weight vector
const FLOAT32 *	p_bias		Bias vector
Const WORD32 * const	P_inp_shape	num_inp_dims	Input shape
WORD32	num_inp_dims	1	Number of dimensions in the input
WORD32	axis	1	Dimension number for which layer normalization will be calculated
FLOAT32	eps	1	Value used in division operations to avoid divide by zero error
Output			
FLOAT32 *	p_out	$\prod\limits_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Output vector
FLOAT32 *	P_mean	$\prod\limits_{i=0}^{i=axis-1} p\_inp\_shape[i]$ (OR) 1	Mean vector
FLOAT32 *	P_rstd		Inverse of standard deviation vector

Returns

0: no error

-1: error, invalid parameters

Restrictions

Arguments	Restrictions
p_inp, p_out, p_weight, p_bias, p_mean, p_rstd, p_inp_shape	Aligned on (size of one element)-byte boundary
	Cannot be NULL

P_out, p_mean, p_std	Must not overlap with inputs
eps, num_inp_dims	Greater than 0
axis	[0, num_inp_dims)

2.4 Reorg Kernels

2.4.1 Slice

Description

The Slice kernels process the input data based on the parameters: axis, start, stop, and stride. The axis parameter indicates the dimension along which the slicing will occur. The operation begins at the position specified by the start parameter and selects elements according to the stride value, continuing until it reaches the stop point within that dimension.

Function variants available are `xa_nn_slice.`

Precision

The kernel is designed to manage all the specified variants with a single implementation. It requires an input parameter, "elm_size" which specifies the size of the data type. By utilizing this parameter, the kernel is capable of supporting all the precision variants listed below.

| Type	Description
8_8	8-bit input, 8-bit output
16_16	16-bit input, 16-bit output
32_32	32-bit input, 32-bit output
8u_8u	Unsigned 8-bit input, 8-bit output
16u_16u	Unsigned 16-bit input, 16-bit output
32u_32u	Unsigned 32-bit input, 32-bit output

Algorithm

bytes_to_copy = inp_shape[axis+1]* inp_shape[axis+2]*...* inp_shape[num_inp_dims-1] * size_of_element
for I = 0 to inp_shape[0]*inp_shape[1]*...*inp_shape[axis-1]
 for J = start to stop insteps of stride
     memcpy(out, inp, bytes_to_copy);
     inp = inp + stride*bytes_to_copy;
     out = out + bytes_to_copy;
 end
end

For ex: For an input vector with dimensions [20][12][4][3][3], and with the following parameters: axis = 2, start = 0, stride = 2, and stop = 3, slices of size 3x3 are extracted according to these parameters. Specifically, slices are taken from the input along the specified axis (axis 2) starting from the index given

by the start parameter (0), selecting elements with the specified stride (2), and ending at the stop parameter (3). This slicing process is performed across the top two dimensions of the input, which are dimensions 20 and 12. This is represented in the below figure.
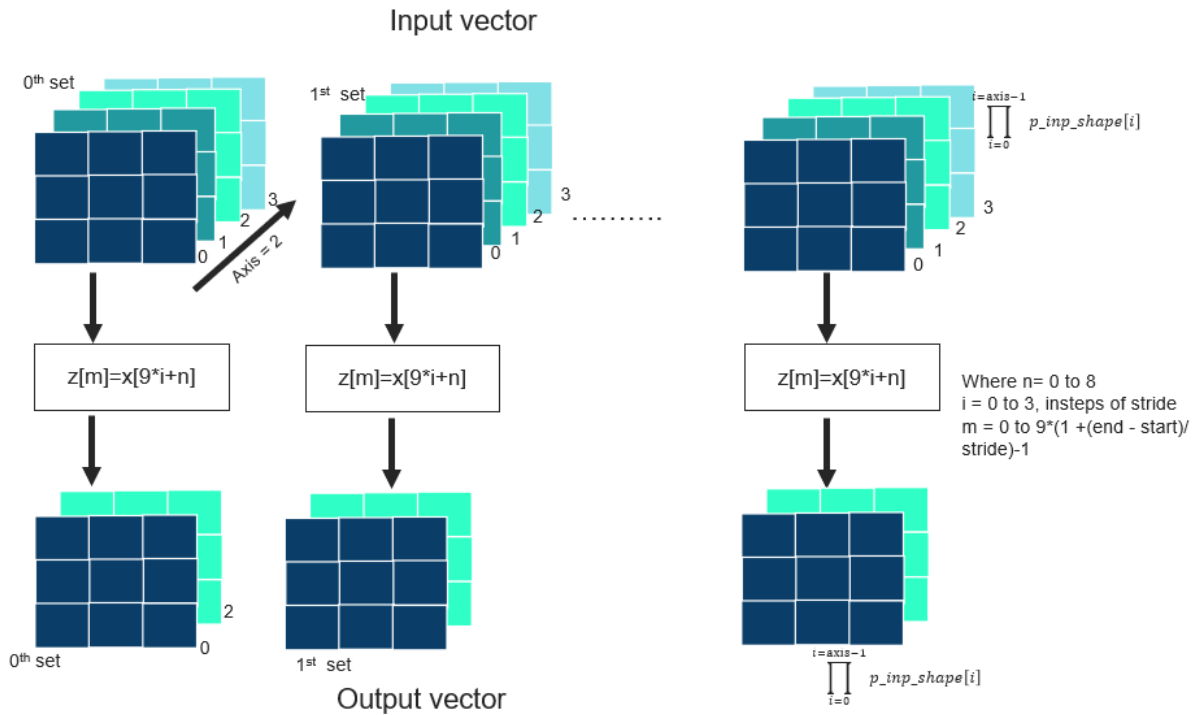
Prototype

```
WORD32 xa_nn_slice(
 WORD8 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
 const   WORD8 * __restrict__  p_inp,
 const WORD32 *const p_inp_shape,
 WORD32 num_inp_dims,
 WORD32 start,
 WORD32 end,
 WORD32 step,
 WORD32 axis,
 WORD32 elm_size);
```

Arguments

Type	Name	Size	Description
Input			
const WORD8 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input vector
Const	p_inp_shape	num_inp_dims	Shape of the input vector

Type	Name	Size	Description
WORD32 *const			
Const WORD32 *const	p_out_shape	num_inp_dims	Shape of the output vector
WORD32	num_inp_dims	1	Number of dimensions in the input vector
WORD32	start	1	Starting index along the axis
WORD32	end	1	Ending index along the axis
WORD32	step	1	stride
WORD32	axis	1	Axis along which the vector has to be sliced
WORD32	elm_size	1	Size of the element. If the tensors of type WORD32, UWORD32 – 4bytes WORD16, UWORD16 – 2 bytes WORD8, UWORD8 – 1 byte
Output			
WORD8 *	p_out	$\prod_{i=0}^{i=axis-1} p\_inp\_shape[i] * ....$ $inp\_shape[1+(end-start)/step]*$ $\prod_{i=axis+1}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Output vector

Returns

0: no error

-1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_inp, p_out, p_inp_shape, p_out_shape	Cannot be NULL Aligned on size of element boundary
p_out	Must not overlap with input
axis	[0, num_inp_dims-1]
step	Greater than 0
start	[0, inp_shape[axis]-1]
end	[0, inp_shape[axis]-1], end>=start
elm_size	1, 2, 4

2.4.2 Permute

Description

This kernel performs a permute operation on a N-dimensional input tensor (up to 5D) as per the combination of dimensions specified in the permute vector. The output tensor's dimension `i` will correspond to the input dimension `permute_vec[i].`

The number of input dimensions must be less than or equal to 5. The 1/2/3/4-dimensional inputs can be scaled up to 5D. The output shape should be conformant with respect to the values in the permute vector.

The naming convention used for the Permute kernel is as follows:

`xa_nn_permute`

Note: Tranpose is a variant of permute operation. So, a tensor transpose can be achieved using the permute kernel by changing the dimension in `permute_vec.`

Precision

The kernel is designed to manage all the specified variants with a single implementation. It requires an input parameter, "elm_size" which specifies the size of the data type. By utilizing this parameter, the kernel is capable of supporting all the precision variants listed below.

Type	Description
8_8	Signed 8-bit input, signed 8-bit output.
16_16	Signed 16-bit input, signed 16-bit output.
32_32	Signed 32-bit input, signed 32-bit output
8u_8u	Unsigned 8-bit input, unsigned 8-bit output.
16u_16u	Unsigned 16-bit input, unsigned 16-bit output.
32u_32u	Unsigned 32-bit input, unsigned 32-bit output

Algorithm

For input P and output Q,
size(Q) = [dim3,dim2,dim4,dim0,dim1] for size(P) = [dim0,dim1,dim2,dim3,dim4] if
permute_vec = [3,2,4,0,1]

For point p in P, and point q in Q,
q(y,x,z,v,w) = p(v,w,x,y,z)
where,

v = 0....dim0 – 1
w = 0....dim1 – 1
x = 0....dim2 – 1
y = 0....dim3 – 1
z = 0....dim4 – 1

Prototype

```
WORD32 xa_nn_permute
(WORD8 * __restrict__ p_out,
 const WORD32 *const p_out_shape,
```

```
const WORD8 * __restrict__ p_inp,
const WORD32 *const p_inp_shape,
const WORD32 * __restrict__ p_permute_vec,
WORD32 num_inp_dims,
WORD32 elm_size);
```

Arguments

Type	Name	Size	Description
Input			
const WORD32 *	p_out_shape	num_inp_dims	Shape of output
const WORD8 *	p_inp	$\prod_{i=0}^{i=num\text{-}inp\text{-}dims-1} p\_inp\_shape[i]$	Input vector. The size of each element is denoted by elm_size parameter.
const WORD32 *	p_inp_shape	num_inp_dims	Shape of input
const WORD32 *	p_permute_vec	num_inp_dims	Permute Vector
WORD32	num_inp_dims	1	Number of input dimensions
WORD32	elm_size	1	Size of the element. If the tensors of type WORD32, UWORD32 – 4bytes WORD16, UWORD16 – 2 bytes WORD8, UWORD8 – 1 byte
Output			
WORD8 *, WORD16 * WORD32 * UWORD8 *, UWORD16 * UWORD32 *	p_out	$\prod_{i=0}^{i=num\text{-}inp\text{-}dims-1} p\_out\_shape[i]$	Output

Returns

- 0: no error

- -1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_out, p_inp	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Must not overlap
p_out_shape, p_inp_shape	Aligned on a 4-byte boundary
	Cannot be NULL
	Must not overlap
	All elements must be greater than zero
p_out_shape	p_out_shape[i] = p_inp_shape[p_permute_vec[i]]
p_permute_vec	Cannot be NULL
num_inp_dims	Must be in the range [1, 5].

elm_size	1, 2, 4

2.4.3 Cat

Description

The Cat kernel concatenates the given inputs into a single output along the dimension specified by the axis parameter. For example, 2 inputs of shapes (1, 8, 128, 32) and (1, 16, 128, 32) are concatenated into an output of shape (1, 24, 128, 32) with axis as '1'.

Function variants available are `xa_nn_cat`

Precision

The kernel is designed to manage all the specified variants with a single implementation. It requires an input parameter, "elm_size" which specifies the size of the data type. By utilizing this parameter, the kernel is capable of supporting all the precision variants listed below.

Type	Description
8_8	8-bit input, 8-bit output
16_16	16-bit input, 16-bit output
32_32	32-bit input, 32-bit output
8u_8u	Unsigned 8-bit input, 8-bit output
16u_16u	Unsigned 16-bit input, 16-bit output
32u_32u	Unsigned 32-bit input, 32-bit output

Algorithm

inp_dims[num_inp][num_dims]
out_dim[num_dims]

For axis = 2
i = 0 to num_inp - 1
    out(d0, d1, sum(inp_dims[0][2] to inp_dims[i-1][2]) + d2, d3, d4, d5) = inp[i](d0, d1, d2, d3, d4, d5)

  d0 = 0 to inp_dims[i][0]
  d1 = 0 to inp_dims[i][1]
  d2 = 0 to inp_dims[i][2]
  d3 = 0 to inp_dims[i][3]
  d4 = 0 to inp_dims[i][4]
  d5 = 0 to inp_dims[i][5]

  if j != axis
    inp_dims[i][j] should be equal to out_dim[j]

  if j == axis
    out_dim[j] == sum(inp_dims[0][j] ... inp_dims[num_inp - 1][j]

Prototype

```
WORD32 xa_nn_cat
  (WORD8 * __restrict__ p_out,   const WORD32 *const p_out_shape
  ,const WORD8 **p_inps,          const WORD32 *const *pp_inps_shape
  ,WORD32 num_inp_dims,           WORD32 num_inp
  ,WORD32 axis,                   WORD32 elm_size);
```

Arguments

Type	Name	Size	Description
Input			
const WORD8 **	pp_inps		Inputs
const WORD32 *	p_out_shape		Shape of output
const WORD32 **	pp_inps_shape		Shape of Inputs
WORD32	num_inp_dims		Number of input dimensions
WORD32	num_inp		Number of Inputs
WORD32	axis		Dimension to concat
WORD8	elm_size		Size of the element. If the tensors of type WORD32, UWORD32 – 4bytes WORD16, UWORD16 – 2 bytes WORD8, UWORD8 – 1 byte
Output			
WORD8 *	p_out		Output

Returns

- 0: no error

- -1: error, invalid parameters

Restrictions:

Arguments	Restrictions
p_out, pp_inps, p_out_shape, pp_inps_shape	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Must not overlap
	The size of each dimension in output except at axis must match with the size of the corresponding each input dimension
p_out	Must not overlap with input
num_inp_dims	Greater than 0
num_inp	Greater than 0
axis	Greater than 0 and less than num_out_dims
elm_size	1, 2, 4

3. Making the library

The Fusion G3 NN library will be released as .tgz file for linux/makefile based usage. The detail about building the library is provided below.

1. Go to directory `libxa_nnlib/build`.
2. From the command prompt, enter:
   ```
   xt-make -f makefile clean all install
   ```

The NN library `xa_nnlib.a` is built and copied to the `lib` directory.

To create a debug build, pass DEBUG=1 makefile option in the make command.